

Week 4 Assignment - Descriptive Modeling

Name: Kesav Adithya Venkidusamy

Course: DSC630 - Predictive Analytics

Instructor: Fadi Alsaleem

You will be using the dataset als_data.csv to apply clustering methods for this assignment. This data gives anonymized data on ALS patients. With this data, complete the following steps:

Importing all the libraries required for this exercise

```
In [26]: ## Importing libraries required for this assignment
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import StandardScaler
from sklearn.cluster import KMeans
from sklearn import metrics
from scipy.spatial.distance import cdist
```

```
In [13]: ## Display all columns in pandas dataframe
pd.set_option('display.max_columns', None)
pd.set_option('display.max_rows', None)
```

Load the Dataset into dataframe

```
In [14]: ## Load the ALS data into a dataframe
als_df = pd.read_csv('als_data.csv')
als_df.head(5)
```

```
Out[14]:
```

	ID	Age_mean	Albumin_max	Albumin_median	Albumin_min	Albumin_range	ALSFRS_slope	ALSFRS
0	1	65	57.0	40.5	38.0	0.066202	-0.965608	
1	2	48	45.0	41.0	39.0	0.010453	-0.921717	
2	3	38	50.0	47.0	45.0	0.008929	-0.914787	
3	4	63	47.0	44.0	41.0	0.012111	-0.598361	
4	5	63	47.0	45.5	42.0	0.008292	-0.444039	

```
In [15]: ## Printing number of rows and columns of als dataframe
als_df.shape
```

```
Out[15]: (2223, 101)
```

```
In [16]: ## Printing the dtype for each of the column
als_df.dtypes
```

```

Out[16]: ID int64
Age_mean int64
Albumin_max float64
Albumin_median float64
Albumin_min float64
Albumin_range float64
ALSFRS_slope float64
ALSFRS_Total_max int64
ALSFRS_Total_median float64
ALSFRS_Total_min int64
ALSFRS_Total_range float64
ALT.SGPT._max float64
ALT.SGPT._median float64
ALT.SGPT._min float64
ALT.SGPT._range float64
AST.SGOT._max int64
AST.SGOT._median float64
AST.SGOT._min float64
AST.SGOT._range float64
Bicarbonate_max float64
Bicarbonate_median float64
Bicarbonate_min float64
Bicarbonate_range float64
Blood.Urea.Nitrogen..BUN._max float64
Blood.Urea.Nitrogen..BUN._median float64
Blood.Urea.Nitrogen..BUN._min float64
Blood.Urea.Nitrogen..BUN._range float64
bp_diastolic_max int64
bp_diastolic_median float64
bp_diastolic_min int64
bp_diastolic_range float64
bp_systolic_max int64
bp_systolic_median float64
bp_systolic_min int64
bp_systolic_range float64
Calcium_max float64
Calcium_median float64
Calcium_min float64
Calcium_range float64
Chloride_max float64
Chloride_median float64
Chloride_min float64
Chloride_range float64
Creatinine_max float64
Creatinine_median float64
Creatinine_min float64
Creatinine_range float64
Gender_mean int64
Glucose_max float64
Glucose_median float64
Glucose_min float64
Glucose_range float64
hands_max int64
hands_median float64
hands_min int64
hands_range float64
Hematocrit_max float64
Hematocrit_median float64
Hematocrit_min float64
Hematocrit_range float64
Hemoglobin_max float64
Hemoglobin_median float64
Hemoglobin_min float64

```

```

Hemoglobin_range      float64
leg_max                int64
leg_median             float64
leg_min               int64
leg_range             float64
mouth_max             int64
mouth_median          float64
mouth_min            int64
mouth_range          float64
onset_delta_mean      int64
onset_site_mean       int64
Platelets_max         int64
Platelets_median      float64
Platelets_min         float64
Potassium_max         float64
Potassium_median      float64
Potassium_min         float64
Potassium_range       float64
pulse_max             int64
pulse_median          float64
pulse_min            int64
pulse_range          float64
respiratory_max       int64
respiratory_median    float64
respiratory_min      int64
respiratory_range     float64
Sodium_max            float64
Sodium_median         float64
Sodium_min           float64
Sodium_range         float64
SubjectID            int64
trunk_max            int64
trunk_median         float64
trunk_min           int64
trunk_range         float64
Urine.Ph_max         float64
Urine.Ph_median      float64
Urine.Ph_min         float64
dtype: object

```

```

In [17]: ## Looking at summary information about your data (total, mean, min, max, freq, unique,
        als_df.describe())

```

```

Out[17]:

```

	ID	Age_mean	Albumin_max	Albumin_median	Albumin_min	Albumin_range	ALSFRS
count	2223.000000	2223.000000	2223.000000	2223.000000	2223.000000	2223.000000	2223.000000
mean	1214.874944	54.550157	47.011134	43.952542	40.766347	0.013779	-0.000000
std	696.678300	11.396546	3.233980	2.654804	3.193087	0.009567	0.000000
min	1.000000	18.000000	37.000000	34.500000	24.000000	0.000000	-4.000000
25%	614.500000	47.000000	45.000000	42.000000	39.000000	0.009042	-1.000000
50%	1213.000000	55.000000	47.000000	44.000000	41.000000	0.012111	-0.000000
75%	1815.500000	63.000000	49.000000	46.000000	43.000000	0.015873	-0.000000
max	2424.000000	81.000000	70.300000	51.100000	49.000000	0.243902	1.000000

1. Remove any data that is not relevant to the patient's ALS condition.

On analyzing all the features present in the data set, I see SubjectID and ID features are irrelevant to patient's ALS condition. So, these variables can be removed from the dataset.

```
In [18]: # Removing Patients ID and SubjectID fields from the dataset
als_df.drop(["SubjectID", "ID"], axis=1, inplace=True)
```

```
In [19]: ## Validating the dataset after removing the columns
als_df.shape
```

```
Out[19]: (2223, 99)
```

```
In [20]: ## Displaying few records in the dataframe
als_df.head()
```

```
Out[20]:
```

	Age_mean	Albumin_max	Albumin_median	Albumin_min	Albumin_range	ALSFRS_slope	ALSFRS_Tot
0	65	57.0	40.5	38.0	0.066202	-0.965608	
1	48	45.0	41.0	39.0	0.010453	-0.921717	
2	38	50.0	47.0	45.0	0.008929	-0.914787	
3	63	47.0	44.0	41.0	0.012111	-0.598361	
4	63	47.0	45.5	42.0	0.008292	-0.444039	

2. Apply a standard scalar to the data.

```
In [22]: ## Define standard scaler
scaler = StandardScaler()

## Transform data
X = scaler.fit_transform(als_df)
print(X)
```

```
[[ 0.91713698  3.08941722 -1.30078105 ... -0.88037551  0.46305355
  1.86853157]
 [-0.57487867 -0.62201561 -1.11240084 ...  0.1926645  -1.13720768
 -0.41915124]
 [-1.45253494  0.92441474  1.14816173 ... -0.88037551 -1.13720768
 -0.41915124]
 ...
 [-0.6626443  -0.31272954  0.01788044 ...  2.33874452  0.46305355
 -0.41915124]
 [-1.54030057  0.61512867  0.01788044 ... -0.88037551 -1.13720768
 -0.41915124]
 [-0.57487867  0.3058426  0.39464087 ... -1.95341552 -1.13720768
 -0.41915124]]
```

```
In [23]: ## Calculating the shape of dataframe
X.shape
```

```
Out[23]: (2223, 99)
```

```
In [25]: ## Calculating the mean and standard deviation of the dataset
         np.mean(X), np.std(X)
```

```
Out[25]: (-8.908541299845311e-17, 1.0)
```

Observation:

1. The mean value is close to zero as expected as the mean value should be close to 0 after applying StandardScalar
2. The Standard deviation is 1 which is also as expected

```
In [27]: ## Reference: https://www.geeksforgeeks.org/elbow-method-for-optimal-value-of-k-in-kmea
         ## Building the clustering model and calculating the values of the Distortion and Inert

         distortions = []
         inertias = []
         mapping1 = {}
         mapping2 = {}
         K = range(1, 10)

         for k in K:
             # Building and fitting the model
             kmeanModel = KMeans(n_clusters=k).fit(X)
             kmeanModel.fit(X)


             distortions.append(sum(np.min(cdist(X, kmeanModel.cluster_centers_, 'euclidean'), ax
             inertias.append(kmeanModel.inertia_)

             mapping1[k] = sum(np.min(cdist(X, kmeanModel.cluster_centers_, 'euclidean'), axis=1)
             mapping2[k] = kmeanModel.inertia_
```

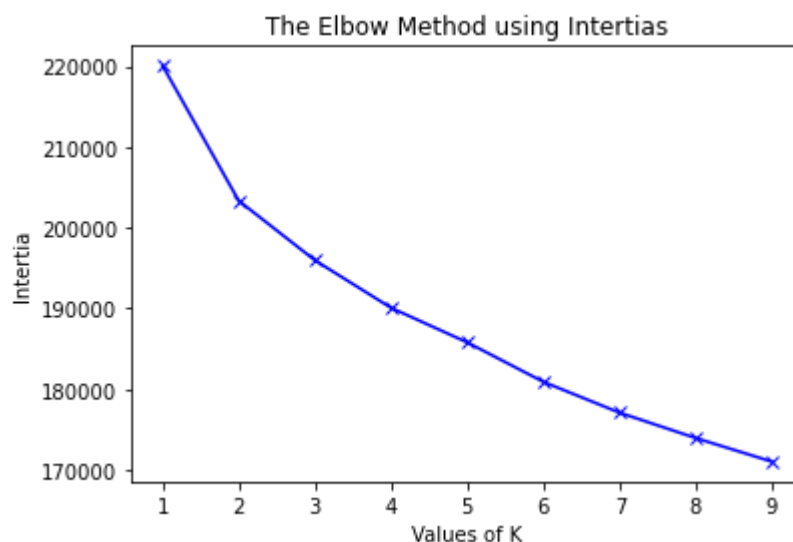
```
In [28]: ## Tabulating and Visualizing the results
         ## Using the different values of Distortion
         for key, val in mapping1.items():
             print(f'{key} : {val}')
```

```
1 : 9.466971249370545
2 : 9.063781809801627
3 : 8.890733728366843
4 : 8.734143328933625
5 : 8.621578780871832
6 : 8.557595851395476
7 : 8.47019365921882
8 : 8.392994844892927
9 : 8.357677447737448
```

```
In [31]: # Plot Elbow method to find K
         plt.plot(K, distortions, 'bx-')
         plt.xlabel('Values of K')
         plt.ylabel('Distortion')
         plt.title('The Elbow Method using Distortion')
         plt.show()
```



```
In [32]: # Plot Elbow method to find K
plt.plot(K, inertias, 'bx-')
plt.xlabel('Values of K')
plt.ylabel('Intertia')
plt.title('The Elbow Method using Intertias')
plt.show()
```



3. Create a plot of the cluster silhouette score versus the number of clusters in a K-means cluster.

```
In [38]: ## Importing the libraries required for silhousette
from sklearn.datasets import make_blobs
from sklearn.cluster import KMeans
from sklearn.metrics import silhouette_samples, silhouette_score
import matplotlib.cm as cm
import matplotlib.style as style
```

```
In [39]: ## Plotting the cluster silhouette score versus the number of clusters in a K-means clu

range_n_clusters = [2, 3, 4, 5, 6, 7, 8, 9, 10]
silhouette_avg_n_clusters = []
```

```

for n_clusters in range_n_clusters:
    # Create a subplot with 1 row and 2 columns
    fig, (ax1, ax2) = plt.subplots(1, 2)
    fig.set_size_inches(18, 7)

    # The 1st subplot is the silhouette plot
    # The silhouette coefficient can range from -1, 1 but in this example all
    # lie within [-0.1, 1]
    ax1.set_xlim([-0.1, 1])

    # The (n_clusters+1)*10 is for inserting blank space between silhouette
    # plots of individual clusters, to demarcate them clearly.
    ax1.set_ylim([0, len(X) + (n_clusters + 1) * 10])

    # Initialize the clusterer with n_clusters value and a random generator
    # seed of 10 for reproducibility.
    clusterer = KMeans(n_clusters=n_clusters, random_state=42)
    cluster_labels = clusterer.fit_predict(X)

    # The silhouette_score gives the average value for all the samples.
    # This gives a perspective into the density and separation of the formed
    # clusters
    silhouette_avg = silhouette_score(X, cluster_labels)
    print("For n_clusters =", n_clusters,
          "The average silhouette_score is :", silhouette_avg)

    silhouette_avg_n_clusters.append(silhouette_avg)
    # Compute the silhouette scores for each sample
    sample_silhouette_values = silhouette_samples(X, cluster_labels)

    y_lower = 10
    for i in range(n_clusters):
        # Aggregate the silhouette scores for samples belonging to
        # cluster i, and sort them
        ith_cluster_silhouette_values = \
            sample_silhouette_values[cluster_labels == i]

        ith_cluster_silhouette_values.sort()

        size_cluster_i = ith_cluster_silhouette_values.shape[0]
        y_upper = y_lower + size_cluster_i

        color = cm.nipy_spectral(float(i) / n_clusters)
        ax1.fill_betweenx(np.arange(y_lower, y_upper),
                          0, ith_cluster_silhouette_values,
                          facecolor=color, edgecolor=color, alpha=0.7)

        # Label the silhouette plots with their cluster numbers at the middle
        ax1.text(-0.05, y_lower + 0.5 * size_cluster_i, str(i))

        # Compute the new y_lower for next plot
        y_lower = y_upper + 10 # 10 for the 0 samples

    ax1.set_title("The silhouette plot for the various clusters.")
    ax1.set_xlabel("The silhouette coefficient values")
    ax1.set_ylabel("Cluster label")

    # The vertical line for average silhouette score of all the values
    ax1.axvline(x=silhouette_avg, color="red", linestyle="--")

```

```

ax1.set_yticks([]) # Clear the yaxis labels / ticks
ax1.set_xticks([-0.1, 0, 0.2, 0.4, 0.6, 0.8, 1])

# 2nd Plot showing the actual clusters formed
colors = cm.nipy_spectral(cluster_labels.astype(float) / n_clusters)
ax2.scatter(X[:, 0], X[:, 1], marker='.', s=30, lw=0, alpha=0.7,
            c=colors, edgecolor='k')

# Labeling the clusters
centers = clusterer.cluster_centers_
# Draw white circles at cluster centers
ax2.scatter(centers[:, 0], centers[:, 1], marker='o',
            c="white", alpha=1, s=200, edgecolor='k')

for i, c in enumerate(centers):
    ax2.scatter(c[0], c[1], marker='.$d$' % i, alpha=1,
                s=50, edgecolor='k')

ax2.set_title("The visualization of the clustered data.")
ax2.set_xlabel("Feature space for the 1st feature")
ax2.set_ylabel("Feature space for the 2nd feature")

plt.suptitle(("Silhouette analysis for KMeans clustering on sample data "
            "with n_clusters = %d" % n_clusters),
            fontsize=14, fontweight='bold')

plt.show()







style.use("fivethirtyeight")
plt.plot(range_n_clusters, silhouette_avg_n_clusters)
plt.xlabel("Number of Clusters (k)")
plt.ylabel("silhouette score")
plt.show()







```





```

For n_clusters = 2 The average silhouette_score is : 0.07878005888570402
For n_clusters = 3 The average silhouette_score is : 0.0687707291658565
For n_clusters = 4 The average silhouette_score is : 0.06973816142698218
For n_clusters = 5 The average silhouette_score is : 0.05697679932842005
For n_clusters = 6 The average silhouette_score is : 0.06477886829610223
For n_clusters = 7 The average silhouette_score is : 0.05187647631845004
For n_clusters = 8 The average silhouette_score is : 0.04954004349267961
For n_clusters = 9 The average silhouette_score is : 0.04393719582297171
For n_clusters = 10 The average silhouette_score is : 0.046121611845315456

```


Silhouette analysis for KMeans clustering on sample data with n_clusters = 2**Silhouette analysis for KMeans clustering on sample data with n_clusters = 3****Silhouette analysis for KMeans clustering on sample data with n_clusters = 4**

Silhouette analysis for KMeans clustering on sample data with n_clusters = 5**Silhouette analysis for KMeans clustering on sample data with n_clusters = 6****Silhouette analysis for KMeans clustering on sample data with n_clusters = 7**

Silhouette analysis for KMeans clustering on sample data with n_clusters = 8**Silhouette analysis for KMeans clustering on sample data with n_clusters = 9****Silhouette analysis for KMeans clustering on sample data with n_clusters = 10**



4. Use the plot created in (3) to choose an optimal number of clusters for K-means. Justify your choice.

Observation

1. None of the silhouettes have clusters that are below-average silhouette scores, so none of the cluster values can be considered bad picks for the given dataset.
2. However, plotting the average silhouette scores for each k shows that the best choice for k is 2 since it has the maximum score (0.07878005888570402). This can also be validated in the Elbow method plot above.

5. Fit a K-means model to the data with the optimal number of clusters chosen in part (4).

```
In [40]: ## Using the cluster value as 2
kmeans = KMeans(init="random", n_clusters=2, n_init=10, max_iter=300, random_state=42)
```

```
In [41]: ## Fit the model to the data with cluster value as 2
kmeans.fit(X)
```

```
Out[41]: KMeans(init='random', n_clusters=2, random_state=42)
```

6. Fit a PCA transformation with two features to the scaled data.

```
In [43]: ## Import PCA Library
from sklearn.decomposition import PCA
```

```
In [44]: ## Reference 1: https://365datascience.com/tutorials/python-tutorials/pca-k-means/
## Reference 2: https://towardsdatascience.com/pca-using-python-scikit-learn-e653f8989e
## Initialize PCA for 2 features
pca_als = PCA(n_components=2)
pca_als.fit(X)
```

```
Out[44]: PCA(n_components=2)
```

```
In [45]:
```

```
# Transform PCA
pca_als.transform(X)
```

```
Out[45]: array([[ -1.4267157 , -2.31976744],
        [ -1.44023996, -4.87137188],
        [  1.6178652 , -0.42935149],
        ...,
        [-0.43290564,  4.2452173 ],
        [-0.330793  ,  3.31718127],
        [  1.46800221,  0.58155086]])
```

```
In [49]: # Creating a variable called score to store transformed PCA
score_pca_als = pca_als.transform(X)
```

```
In [51]: # Create a new dataframe with the original features and add the PCA scores and assigned
df_segmn_pca_kmeans = pd.concat([als_df.reset_index(drop=True), pd.DataFrame(score_pca_als)], axis=1)
df_segmn_pca_kmeans.columns.values[-2:] = ['PCA_1', 'PCA_2']
```

```
In [52]: # The last column we add contains the pca k-means clustering labels
df_segmn_pca_kmeans['Segment K-means PCA'] = kmeans.labels_
```

```
In [56]: # Add the names of the segments to the labels
df_segmn_pca_kmeans['Segment'] = df_segmn_pca_kmeans['Segment K-means PCA'].map({0:'first', 1:'second', 2:'third', 3:'fourth', 4:'fifth'})
```

```
In [57]: ## Show few records from the dataframe
df_segmn_pca_kmeans.head()
```

```
Out[57]:
```

	Age_mean	Albumin_max	Albumin_median	Albumin_min	Albumin_range	ALSFRS_slope	ALSFRS_Tot
0	65	57.0	40.5	38.0	0.066202	-0.965608	
1	48	45.0	41.0	39.0	0.010453	-0.921717	
2	38	50.0	47.0	45.0	0.008929	-0.914787	
3	63	47.0	44.0	41.0	0.012111	-0.598361	
4	63	47.0	45.5	42.0	0.008292	-0.444039	


7. Make a scatterplot of the PCA transformed data coloring each point by its cluster value.

```
In [58]: # Plot scatterplot of PCA transformed data coloring each point by its cluster value
x_axis = df_segmn_pca_kmeans['PCA_2']
y_axis = df_segmn_pca_kmeans['PCA_1']
plt.figure(figsize=(10,8))
sns.scatterplot(x_axis,y_axis, hue=df_segmn_pca_kmeans['Segment'], palette = ['g', 'r'])
plt.title('Clusters by PCA Components')
plt.show()
```

C:\Users\KesavAdithya\anaconda3\lib\site-packages\seaborn_decorators.py:36: FutureWarning

ng: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

warnings.warn(



In [61]:

```
## Calculating Eigen values for each components
print('Eigenvalues: {}'.format(pca_als.explained_variance_))
print('Explained Variance Ratio: {}'.format(pca_als.explained_variance_ratio_))
```

```
Eigenvalues: [11.22229079  6.38502006]
Explained Variance Ratio: [0.11330548 0.06446614]
```

8. Summarize your results and make a conclusion.

1. By plotting the results of K-means algorithm with PCA, we see all the cluster clusters are jumbled all together. However, when we employ PCA prior to using K-means we can visually separate almost the entire data set. That was one of the biggest goals of PCA - to reduce the number of variables by combining them into bigger, more meaningful features.
2. There is some overlap between the red and green segments. But, as a whole, both the segments are clearly separated.
3. Here we can see, the first component explained 11.33% variance and the second component explained 6.4% variance
4. Looking at Eigenvalues, we see Principal Component 1 is ~64% (11.22/17.7) and PC2 is about ~36% (6.3/17.6). This means that the Y axis accounts for ~64% of the variation in the dataset, and X axis accounts for the remaining ~36%.

In []: